

Media Engineering and Technology Faculty
German University in Cairo



LLM Hallucination Detection via Knowledge Augmentation

Bachelor Thesis

Author: Mohammed Mahmoud Elnaggar
Supervisors: Dr. Nourhan Ehab
Submission Date: XX July, 20XX

This is to certify that:

- (i) the thesis comprises only my original work toward the Bachelor Degree
- (ii) due acknowledgement has been made in the text to all other material used

Mohammed Mahmoud Elnaggar
XX July, 20XX

Acknowledgments

Text

Abstract

Abstact

Contents

Acknowledgments	V
1 Introduction	1
1.1 Section Name	1
1.2 Another Section	1
2 Background and Related Works	3
2.1 Understanding hallucinations in Large Language Models	3
2.1.1 Hallucinations definition in LLM	3
2.1.2 LLM hallucinations types and categories	4
2.1.3 Causes of LLM Hallucinations from Training Data	4
2.1.4 Hallucination detection benchmarks	5
2.2 Knowledge Augmentation Approaches	6
2.2.1 Retrieval-Augmented Generation	6
2.2.1.1 Components	7
2.2.1.2 Models	7
2.2.2 Knowledge Graphs as Structured Knowledge Repositories	7
2.2.2.1 Definition and Representational Framework	7
2.2.2.2 Integration with Large Language Models	8
2.3 Controlling and Constraining LLM Output	8
2.3.1 White-Box Approaches: Leveraging Logit Access	8
2.3.2 Black-Box Approaches: Constrained Decoding Without Logits	9
2.4 Black-Box LLM Hallucination Detection	10
2.4.1 Zero-Resource Knowledge Hallucination Detection	10
2.4.2 External Knowledge-Based Hallucination Detection	11
3 Conclusion	13
4 Future Work	15
Appendix	16
A Lists	17
List of Abbreviations	17
List of Figures	18

Chapter 1

Introduction

1.1 Section Name

Some sample text with an Acronym Without Citation (AC), some citation ?, and some more Acronym With Citation ? (AC2).

1.2 Another Section

Reference to Section 1.1, and reuse of AC nad AC2 with also full use of Acronym With Citation ? (AC2).

Chapter 2

Background and Related Works

This chapter explains the foundations needed to understand the approach used in this thesis: a knowledge-grounded verification layer that detects hallucinations in LLM outputs by checking them against explicit symbolic knowledge. We commence with Section 2.1, where we define hallucinations in LLMs, discuss their types and causes, and review the benchmarks used to evaluate detection approaches. Subsequently, in Section 2.2, we explore Retrieval-Augmented Generation (RAG) as a knowledge augmentation architecture, which motivates our approach where we will replace retrieval with symbolic constraint checking, and we introduce knowledge graphs as the structured knowledge representation that will serve as the backbone of our proposed approach. Following this, Section 2.3 examines how we can control LLM outputs, particularly in black-box settings where control typically occurs after generation, as it has several connections to black-box hallucination detection. We include white-box approaches in this section primarily to compare the different methodologies and understand how they differ. In Section 2.4, we finally discuss papers that present approaches for detecting hallucinations in black-box LLMs, covering both zero-resource and external knowledge-based methods.

2.1 Understanding hallucinations in Large Language Models

This section covers the foundational concepts of LLM hallucinations, including their definition, types, causes, and the benchmarks used to evaluate detection approaches.

2.1.1 Hallucinations definition in LLM

In the current era, while we are in the evolution in LLMs that generate fluent responses across various tasks, there is a critical problem that may have a significant impact on certain crucial fields where the information provided in the responses of different LLMs

must be factual. Alongside the evolution of LLMs and their fluent responses, LLMs remain prone to hallucinations. LLM hallucination refers to the phenomenon where an LLM produces responses that are unfaithful or nonsensical relative to the source content [Ji et al., 2023, Huang et al., 2025]. Furthermore, hallucinations have become increasingly difficult to detect, as LLMs produce fluent responses, which may lead to harmful consequences.

2.1.2 LLM hallucinations types and categories

First, there are two traditional types of hallucinations: intrinsic and extrinsic [Ji et al., 2023]. Intrinsic hallucination means that the generated output contradicts the source content — for example, when summarizing a text and contradicting a date mentioned in the original content. Extrinsic hallucination means that the generated output cannot be verified against the source content or an external knowledge base. Second, in more recent papers, there is a further categorization of LLM hallucination types into factuality hallucinations and faithfulness hallucinations, each of which is further subdivided [Huang et al., 2025]. For factuality hallucinations, there are two types: the first is factual contradiction, where the LLM generates output that contradicts facts grounded in reality, and the second is factual fabrication, which refers to generated output that cannot be verified against any known source. As for faithfulness hallucinations, they are categorized into three types. The first is instruction inconsistency, where the LLM output deviates from the user’s given directions. The second is context inconsistency, where the LLM output is unfaithful to the contextual information provided by the user. The third is logical inconsistency, where the LLM exhibits internal logical contradictions within its own output.

2.1.3 Causes of LLM Hallucinations from Training Data

In this section, we explore the root cause of hallucinations in LLMs from data [Huang et al., 2025, Ji et al., 2023].

The data used for training LLMs comprises two primary categories: (1) pre-training data, through which LLMs acquire most of their general knowledge and factual information, and (2) alignment data, which teaches LLMs to align with human preferences [Huang et al., 2025].

Misinformation and Biases Fake news and unfounded data have spread widely across the internet, comprising a significant portion of LLM training data and consequently causing hallucinations. Additionally, certain biases related to hallucination have been observed, particularly those associated with gender and nationality [Huang et al., 2025].

Knowledge Boundaries Although LLMs are equipped with extensive factual knowledge, they inherently possess knowledge boundaries. The long-tail boundary arises from the non-uniform distribution of training data, which causes LLMs to demonstrate varying levels of proficiency across different domains. The up-to-date problem occurs because once an LLM is trained, its knowledge becomes outdated, and factual information that changes over time may lead to hallucinations [Huang et al., 2025].

Copyright-Sensitive Knowledge Since LLMs rely on publicly available data, limitations arise in domains that contain inaccessible or restricted data, which causes hallucinations in those specific domains [Huang et al., 2025].

Inferior Alignment LLMs struggle to acquire new knowledge during Supervised Fine-Tuning (SFT), and a correlation has been identified between the knowledge acquired during SFT and the occurrence of hallucinations [Huang et al., 2025].

2.1.4 Hallucination detection benchmarks

Previous benchmarks were focused on models with specific tasks and were not general purpose. The benchmarks were evaluating smaller and task-specific models which produce errors which are different from new era LLMs, since benchmarks used were on smaller task-specific models which is not compatible with today’s LLMs and cannot be used to evaluate how well detection methods used on today’s powerful LLMs; these benchmarks must be changed [Huang et al., 2025].

A set of benchmarks has since been introduced to evaluate how well hallucination detection methods work, in this part we discuss some of the modern datasets to evaluate hallucination detection approaches.

HaluEval: A large-scale benchmark introduced by [Li et al., 2023] which is a large collection of 35,000 generated and human-annotated samples to evaluate LLMs’ ability to recognize hallucinations. The benchmark covers 30,000 samples across multiple tasks including question answering, knowledge-grounded dialogue, and text summarization, providing a comprehensive evaluation framework for hallucination detection across different LLM applications, along with 5,000 samples of general user queries with ChatGPT responses.

FELM: A benchmark introduced by [Chen et al., 2023] for evaluating factuality evaluation methods themselves, rather than directly evaluating LLMs. It contains 2,000+ responses from various LLMs across different domains, each annotated with fine-grained factuality labels, enabling systematic assessment of how well different evaluation metrics and detectors perform at identifying factual errors.

SelfCheckGPT-Wikibio: The WikiBio GPT-3 dataset introduced by [Manakul et al., 2023] offers a sentence-level dataset created by GPT-3 on individuals from the WikiBio dataset, with manual annotation of factuality. WikiBio is a dataset where each input is the first paragraph of Wikipedia articles of specific concepts. They sampled 238 of the top 20 percent of longest articles.

BAMBOO: A benchmark proposed by [Dong et al., 2023] which is a multi-task long-context benchmark used to evaluate long-text tasks including hallucination detection. BAMBOO consists of 10 different datasets from 5 different long-text understanding tasks: question answering, hallucination detection, text sorting, language modeling, and code completion.

ScreenEval: A benchmark introduced by [Lattimer et al., 2023] for evaluating factual inconsistencies in long dialogues. This dataset comprises 52 dialogues averaging 6,000 tokens per dialogue. The dataset uses TV scripts and summaries pulled from the SummScreen dataset; they also generate summaries using Longformer and GPT-4 on 52 scripts from the SummScreen test set.

PHD: A passage-level hallucination detection benchmark introduced by [Yang et al., 2023] containing 300 samples generated by ChatGPT on Wikipedia articles, with human annotations at both sentence and passage levels. The paper also introduces a reverse validation method that reconstructs questions from generated passages and verifies consistency, specifically designed to address the omission problem in traditional consistency-based detection.

LSum: A benchmark introduced by [Feng et al., 2023] for evaluating hallucination in text summarization tasks. The paper proposes an adversarial approach that decouples comprehension and embellishment abilities of LLMs, demonstrating that factual consistency can be improved by separating what the model understands from what it generates, with applications for both detection and mitigation of hallucinations in summarization.

2.2 Knowledge Augmentation Approaches

2.2.1 Retrieval-Augmented Generation

Pre-trained language models have been shown to learn a substantial amount of in-depth knowledge from data, and they are capable of doing so without accessing any external memory. However, such models cannot expand or revise their knowledge after training, which may lead to hallucinations [Lewis et al., 2020]. This thesis draws on the same motivation as RAG — grounding LLM outputs in external knowledge — but replaces retrieval with symbolic constraint checking against a structured knowledge representation.

2.2.1.1 Components

RAG consists of two core components: the Retriever and the Generator [Lewis et al., 2020].

- **Retriever:** Responsible for returning the top- k documents based on a truncated distribution over the input sequence X .
- **Generator:** Utilizes the retrieved documents as additional context to generate the target sequence Y .

2.2.1.2 Models

- **RAG Sequence Model:** The RAG sequence model uses the same set of retrieved documents to generate the entire target sequence [Lewis et al., 2020].
- **RAG Token Model:** The RAG token model uses different retrieved documents for each individual target token, allowing for more fine-grained document selection during generation [Lewis et al., 2020].

2.2.2 Knowledge Graphs as Structured Knowledge Repositories

2.2.2.1 Definition and Representational Framework

Knowledge graphs represent information as entities (nodes) and relations (edges), forming a structured repository of facts. Facts are stored as triples of the form $\langle \textit{subject}, \textit{predicate}, \textit{object} \rangle$, such as $\langle \textit{Paris}, \textit{capitalOf}, \textit{France} \rangle$, which allows structured reasoning over the stored knowledge.

To illustrate, consider the sentence "Paris is the capital of France." This would be stored as the triple (Paris, capitalOf, France). If an LLM generates the response "Paris is the capital of Germany," a verification system can query the knowledge graph for (Paris, capitalOf, ?). The graph returns "France," which contradicts the generated "Germany," signaling a hallucination. This query-based mechanism is how knowledge graphs enable fact-level verification.

In the context of hallucination detection, knowledge graphs are used as a way of storing rules or facts that will be verified or assigned a score [Sawczyn et al., 2025, Kale and Alfeo, 2025, Fang et al., 2024].

2.2.2.2 Integration with Large Language Models

Recent works integrate knowledge graphs with LLMs to improve factual accuracy and detect hallucinations. The integration typically follows two patterns. The first pattern constructs graphs from LLM responses and uses them for structural consistency checking. For instance, [Fang et al., 2024] constructs graphs from responses and uses graph neural networks to model contextual dependencies between triples, enabling detection of logical inconsistencies across multiple facts. The second pattern uses knowledge graphs as an external reference to verify extracted facts. [Sawczyn et al., 2025] and [Kale and Alfeo, 2025] leverage knowledge graphs for fine-grained fact-level consistency checking by comparing triples extracted from LLM outputs against triples stored in the graph.

2.3 Controlling and Constraining LLM Output

Enforcing domain rules and logical constraints on LLM output is a major concern in modern research, as it makes LLMs more reliable. While this thesis focuses on detecting hallucinations after generation, examining methods that prevent hallucinations through controlled generation provides valuable context. Both control and detection share a common goal: ensuring LLM outputs align with expected facts or rules. Additionally, some detection approaches borrow techniques from constrained decoding, particularly in how they interact with external knowledge sources. There are several approaches for controlling LLM output: ones that are used with white-box LLMs and ones that are used with black-box LLMs, where access to their logits is unavailable. In this section we discuss each one of them, with emphasis on black-box methods as they are most relevant to the detection approaches surveyed in Section 2.4.

2.3.1 White-Box Approaches: Leveraging Logit Access

The techniques used to control LLM output share the same goal: enforcing $p(x|a)$, where x is the output sequence and a is the constraint that needs to be satisfied.

There are papers that distill a TPM (Tractable Probabilistic Model), most commonly an HMM (Hidden Markov Model), then combine the HMM with the LLM’s output probability distribution and search ahead to determine whether the given constraint will be enforced. One such paper, GeLaTo, trains an HMM via maximum likelihood estimation on samples drawn from LLM output, and at generation time combines the HMM probability with the LLM output [Zhang et al., 2023a]. Another paper, Ctrl-G, generalizes GeLaTo by distilling an HMM, constructing a DFA, and then conditioning the HMM on the DFA [Lu et al., 2023]. Yet another paper distills an HMM and combines it with an efficiently trained classifier that guides the computation of the Expected Attribute Probability, which in turn steers generation [Wang et al., 2022].

Another approach is presented in the paper FUDGE, which shares the same goal of modeling $p(x|a)$, but instead of using a TPM, it trains a binary predictor on whether

constraint a will be satisfied based on an incomplete sequence prefix [Yang and Klein, 2021]. It multiplies the probability distributions of the predictor and the LLM output, then selects the token with the highest weight.

These white-box methods share a common pattern: they all require access to the LLM’s internal logits and involve training auxiliary models (HMMs or binary predictors) to guide generation. While they are effective, their reliance on logit access limits their work with black-box LLMs.

2.3.2 Black-Box Approaches: Constrained Decoding Without Logits

In the previous section, constraints were applied by leveraging access to the LLM’s probability distribution. However, when using black-box LLMs such as ChatGPT or Claude, access to the model’s internal logits is unavailable, so alternative approaches are needed to enforce rules or logical constraints on the output of such black-box LLMs.

One such technique is presented in the paper SketchGCD, which proposes an approach for constrained decoding on LLM output that requires no further training [Lu et al., 2024]. It uses a lightweight open-source model to refine the output of a heavyweight black-box LLM to satisfy specific constraints. This approach is divided into two components: a sketcher and a constrained decoder. For the constraints to be applied, it restricts entities and relations to the Wikidata knowledge graph and enforces a structural format that the output must follow.

Another technique that integrates LLMs with symbolic solvers to improve logical problem-solving is LOGIC-LM [Pan et al., 2023], which takes a problem P and a goal G and feeds them into a multi-stage process to derive the answer. First, it prompts an LLM to translate P and G into task-specific symbolic representations, using four different symbolic formulations: logic programming, first-order logic, constraint satisfaction problems, and analytical reasoning. After converting P and G into symbolic representations, they are passed to a symbolic solver to obtain answer A by enforcing logical rules. The approach also includes a self-refinement module: if the solver returns a logical or syntactic error, the LLM revises the formulation and repeats the process iteratively.

These black-box control methods demonstrate that logical and factual constraints can be enforced on LLM outputs without internal logit access. Notably, approaches like SketchGCD rely on knowledge graphs (Wikidata) as the source of constraints—the same type of structured knowledge that this thesis uses for detection. The key difference is one of timing: control methods intervene *during* generation to prevent hallucinations, while detection methods operate *after* generation to identify them. This distinction aside, the underlying mechanism of comparing LLM content against a knowledge graph is common to both, making constrained decoding a relevant adjacent field to the hallucination detection approaches discussed next.

2.4 Black-Box LLM Hallucination Detection

As mentioned previously, hallucinations are a major concern in the LLM field, so detecting them is a crucial area of research in which several approaches have been proposed, each employing different techniques. The hallucination detection approaches for black-box LLMs are divided into two parts: zero-resource and external knowledge hallucination detection.

2.4.1 Zero-Resource Knowledge Hallucination Detection

This section covers four different types of zero-resource approaches: 1) SelfCheck zero-resource hallucination detection, 2) fine-grained cross-model evaluation, 3) graph-based context-aware hallucination detection and 4) knowledge graphs for hallucination self-detection. These methods share a common strength: they require no external knowledge base, making them broadly applicable across domains. However, they also share a fundamental limitation that motivates the external knowledge-based approaches discussed in Section 2.4.2.

SelfCheckGPT: One approach is presented in the paper SelfCheckGPT, which requires no external knowledge and uses the LLM to self-check its output, operating effectively with black-box LLMs [Manakul et al., 2023]. This is achieved by obtaining a response R from the LLM given a prompt, then sampling additional responses, which are subsequently used with different methods (BERTScore, NLI, n-grams, etc.) to measure the consistency between the original response and the samples, and then assigning a hallucination score.

FINCH-ZK: Another approach is presented in the paper FINCH-ZK (FINE-grained Cross-model consistency for Hallucination detection and mitigation with Zero Knowledge) [Karatopak et al., 2024]. It requires no external knowledge base to detect hallucinations and uses an LLM to judge the outputs. First, a cross-model sample generation step is performed, in which variants of the prompt are generated and a set of sampler LLMs (M) are used to produce samples. Once these samples are collected, a judge LLM performs fine-grained evaluation. The response is segmented into blocks, and then a judge LLM (J) is used to evaluate each block against each sample, after which hallucination scores are computed.

Graph-based Context-Aware (GCA): The third approach is Graph-based Context-Aware (GCA) zero-resource hallucination detection, which focuses on long text generation with black-box models [Fang et al., 2024]. The paper begins by extracting knowledge triples from the response by providing the LLM with the response and prompting it to extract the triples. Next, it asks the LLM to verify each triple against the original response, and if there is ambiguity in equivalence, it asks the LLM to adjust. Then they construct a graph for the original response and each sampled response, using a Relational Graph Convolutional Network (RGCN) which passes information through nodes, so that

when the model checks a triple, it remembers information from other triples and can detect logical errors. After that, they perform graph-based consistency comparison where they compare triples from the original response with triples from sampled responses. If alignment is found, they add 1 to the consistency score, and if there is contradiction between triples, they subtract 1 from the consistency score. Finally, they perform reverse verification through triple reconstruction, where they check if the LLM can correctly reconstruct the head entity, relation, and tail entity of each triple using three different tasks, and combine these scores with the consistency score to make the final hallucination judgment.

Knowledge Graph Self-Detection: The fourth approach involves using knowledge graphs for self-detection of hallucinations. One method is presented in the paper Fact-SelfCheck, which operates at the fact level rather than the sentence level [Sawczyn et al., 2025]. It follows a similar procedure to SelfCheckGPT by first generating a response and multiple samples. However, instead of comparing sentences, it constructs knowledge graphs by extracting entities and relations from the response using LLM prompts, does the same with the samples, and then checks the consistency of the extracted facts to assign hallucination scores. Similarly, another approach is presented in the paper Lie to Me, which enhances both single-sample and multi-sample self-detection techniques using knowledge graphs [Kale and Alfeo, 2025]. After obtaining a response, it prompts the LLM to construct a knowledge graph, and then uses either self-questioning or self-confidence approaches for single-sample detection, or integrates knowledge graphs into SelfCheckGPT for multi-sample detection, to assign hallucination scores at the fact level.

While these zero-resource approaches offer practical advantages, they share a critical limitation: they cannot verify claims that fall outside the LLM’s parametric knowledge. If a model confidently generates a false fact and also repeats it in sampled responses, consistency-based methods may incorrectly label it as factual. This limitation necessitates approaches that consult external knowledge sources, which are discussed next.

2.4.2 External Knowledge-Based Hallucination Detection

While zero-resource approaches offer practical advantages, they cannot verify claims that fall outside their parametric knowledge, which is where external knowledge-based methods are needed.

FactScore: One approach is proposed in the paper FActScore, which evaluates generation based on atomic facts [Min et al., 2023]. This approach uses both human annotation and automated annotation to evaluate atomic facts. For the human annotation, they sampled 183 entities from Wikipedia, fed prompts to LLMs, and then human annotators broke the generation into atomic facts and labeled them as supported, not-supported, or irrelevant based on the knowledge source of Wikipedia. However, this approach is expensive, so they introduced automated annotation where an LLM subject generates text and an LLM evaluator assigns fact scores. The estimator breaks the generation into atomic facts and assigns scores against the knowledge source.

REFCHECKER: Another approach is REFCHECKER, which uses claim-triplets instead of sentence-level evaluation, as one sentence may contain several facts [Hu et al., 2024]. This approach consists of two parts: the extractor and the checker. The extractor extracts triples from the response using models like GPT-4, Mixtral, or a fine-tuned Mistral model trained on 10k responses with Mixtral as a teacher. The checker then compares these triples against a reference to determine if they are entailed, contradicted, or neutral.

Both approaches demonstrate the value of consulting external knowledge for hallucination detection. They also share a common characteristic: they verify against unstructured text (Wikipedia), which requires natural language understanding for matching. This thesis takes a different direction: using a structured knowledge graph as the reference, which enables direct triple-to-triple matching without the ambiguity of natural language interpretation. This distinction is explored in the proposed approach presented in the next chapter.

Chapter 3

Conclusion

Conclusion

Chapter 4

Future Work

Text

Appendix

Appendix A

Lists

AC	Acronym Without Citation
AC2	Acronym With Citation ?

List of Figures

Bibliography

- Huang, L., Yu, W., Ma, W., Zhong, W., Feng, Z., Wang, H., Chen, Q., Peng, W., Feng, X., Qin, B., & Liu, T. (2025). A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions. *ACM Transactions on Information Systems*, 43(2), 1–56. <https://doi.org/10.1145/3703155>
- Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., Ishii, E., Bang, Y. J., Madotto, A., & Fung, P. (2023). Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12), 1–38. <https://doi.org/10.1145/3571730>
- Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W.-t., Rocktäschel, T., Riedel, S., & Kiela, D. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems*, 33, 9459–9474. <https://proceedings.neurips.cc/paper/2020/hash/6b493230205f780e1bc26945df7481e5-Abstract.html>
- Manakul, P., Liusie, A., & Gales, M. (2023). SelfCheckGPT: Zero-resource black-box hallucination detection for generative large language models. In H. Bouamor, J. Pino, & K. Bali (Eds.), *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing* (pp. 9004–9017). Association for Computational Linguistics. <https://aclanthology.org/2023.emnlp-main.557/>
- Lu, X., Welleck, S., Hessel, J., Jiang, L., Qin, L., West, P., Ammanabrolu, P., & Choi, Y. (2023). Adaptable logical control for large language models. *arXiv preprint arXiv:2306.00092*. <https://arxiv.org/abs/2306.00092>
- Sawczyn, A., Binkowski, J., Janiak, D., Gabrys, B., & Kajdanowicz, T. (2025). Fact-SelfCheck: Fact-level black-box hallucination detection for LLMs. *arXiv preprint arXiv:2501.07872*. <https://arxiv.org/abs/2501.07872>
- Yang, K., & Klein, D. (2021). FUDGE: Controlled text generation with future discriminators. In K. Toutanova, A. Rumshisky, L. Zettlemoyer, D. Hakkani-Tur, I. Beltagy, S. Bethard, R. Cotterell, T. Chakraborty, & Y. Zhou (Eds.), *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies* (pp. 3511–3535). Association for Computational Linguistics. <https://aclanthology.org/2021.naacl-main.276/>

- Kale, S., & Alfeo, A. L. (2025). Lie to me: Knowledge graphs for robust hallucination self-detection in LLMs. *arXiv preprint arXiv:2509.12345*. <https://arxiv.org/abs/2509.12345>
- Wang, D., Peng, B., Liu, Z., Zhang, N., & Chen, J. (2022). TRACE back from the future: A probabilistic reasoning approach to controllable language generation. *arXiv preprint arXiv:2210.06647*. <https://arxiv.org/abs/2210.06647>
- Zhang, H., Song, H., Li, S., Zhou, M., & Song, D. (2023). Tractable control for autoregressive language generation. *arXiv preprint arXiv:2304.07438*. <https://arxiv.org/abs/2304.07438>
- Karatopak, E., Cekinel, R. F., & Yuret, D. (2024). Zero-knowledge LLM hallucination detection and mitigation through fine-grained cross-model consistency. *arXiv preprint arXiv:2405.16853*. <https://arxiv.org/abs/2405.16853>
- Fang, X., Huang, Z., Tian, Z., Fang, M., Pan, Z., Fang, Q., Wen, Z., Pan, H., & Li, D. (2024). Zero-resource hallucination detection for text generation via graph-based contextual knowledge triples modeling. *arXiv preprint arXiv:2409.11283*. <https://arxiv.org/abs/2409.11283>
- Lu, X., Welleck, S., Hessel, J., Jiang, L., Qin, L., West, P., Ammanabrolu, P., & Choi, Y. (2024). Sketch-guided constrained decoding for boosting blackbox large language models without logit access. *arXiv preprint arXiv:2401.09967*. <https://arxiv.org/abs/2401.09967>
- Min, S., Krishna, K., Lyu, X., Lewis, M., Yih, W., Koh, P. W., Iyyer, M., Zettlemoyer, L., & Hajishirzi, H. (2023). FActScore: Fine-grained Atomic Evaluation of Factual Precision in Long Form Text Generation. In H. Bouamor, J. Pino, & K. Bali (Eds.), *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing* (pp. 12076–12100). Association for Computational Linguistics. <https://doi.org/10.18653/v1/2023.emnlp-main.741>
- Hu, X., Ru, D., Qiu, L., Guo, Q., Zhang, T., Xu, Y., Luo, Y., Liu, P., Zhang, Y., & Zhang, Z. (2024). Knowledge-Centric Hallucination Detection. In Y. Al-Onaizan, M. Bansal, & Y. Chen (Eds.), *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing* (pp. 6953–6975). Association for Computational Linguistics. <https://doi.org/10.18653/v1/2024.emnlp-main.395>
- Li, J., Cheng, X., Zhao, W. X., Nie, J.-Y., & Wen, J.-R. (2023). HaluEval: A Large-Scale Hallucination Evaluation Benchmark for Large Language Models. *CoRR*, abs/2305.11747. <https://doi.org/10.48550/arXiv.2305.11747>
- Dong, Z., Tang, T., Li, J., Zhao, W. X., & Wen, J.-R. (2023). BAMBOO: A Comprehensive Benchmark for Evaluating Long Text Modeling Capacities of Large Language Models. *arXiv preprint*, abs/2309.13345. <https://arxiv.org/abs/2309.13345>

- Lattimer, B. M., Chen, P., Zhang, X., & Yang, Y. (2023). Fast and Accurate Factual Inconsistency Detection Over Long Documents. *arXiv preprint*, abs/2310.13189. <https://arxiv.org/abs/2310.13189>
- Chen, S., Zhao, Y., Zhang, J., Chern, I.-C., Gao, S., Liu, P., & He, J. (2023). FELM: Benchmarking Factuality Evaluation of Large Language Models. *arXiv preprint*, abs/2310.00741. <https://arxiv.org/abs/2310.00741>
- Yang, S., Sun, R., & Wan, X. (2023). A New Benchmark and Reverse Validation Method for Passage-level Hallucination Detection. *arXiv preprint*, abs/2310.06498. <https://arxiv.org/abs/2310.06498>
- Feng, H., Fan, Y., Liu, X., Lin, T.-E., Yao, Z., Wu, Y., Huang, F., Li, Y., & Ma, Q. (2023). Improving Factual Consistency of Text Summarization by Adversarially Decoupling Comprehension and Embellishment Abilities of LLMs. *CoRR*, abs/2310.19347. <https://doi.org/10.48550/arXiv.2310.19347>
- Pan, L., Albalak, A., Wang, X., & Wang, W. Y. (2023). Logic-LM: Empowering large language models with symbolic solvers for faithful logical reasoning. *arXiv preprint*, abs/2305.12295. <https://arxiv.org/abs/2305.12295>